

Isolation Heuristics Analysis

Artificial Intelligence Nanodegree

February 16, 2017

Terence So

Summary

Three custom players were evaluated for this project: **Aggressive**, **Balanced**, and **Monte Carlo**, each representing the primary strategy of the custom player in game of Isolation. Each player was evaluated using a tweaked version of the provided tournament format against a baseline strategy **ID_Improved** which has a winrate of **60.00%**. The following are the results in summary:

- The Aggressive strategy performed near ID_Improved with a winrate of **63.57%**.
- The Balanced Strategy performed no better with a **63.21%** winrate.
- The Monte Carlo strategy had the best performance with a winrate of **74.29%**.

Measuring Heuristic Performance

During the course of this project, it was observed that results varied from simulation to simulation. For example, ID_Improved has been observed to have winrates from 55% to about 63% with mixed match results with other players. This holds true with each evaluated strategy. In an effort to increase the accuracy of final score, the tournament format was slightly modified by bumping up the number of games played per match from 20 to 40.

Strategy Evaluation

ID_Improved

Formula: `player moves - opponent moves`

This is the baseline player or the player 'to beat'. The results of one its tournament runs is as follows:

```
*****
Evaluating: ID_Improved
*****

Playing Matches:
-----
Match 1: ID_Improved vs Random      Result: 37 to 3
Match 2: ID_Improved vs MM_Null     Result: 27 to 13
Match 3: ID_Improved vs MM_Open     Result: 20 to 20
Match 4: ID_Improved vs MM_Improved Result: 18 to 22
Match 5: ID_Improved vs AB_Null     Result: 23 to 17
Match 6: ID_Improved vs AB_Open     Result: 22 to 18
Match 7: ID_Improved vs AB_Improved Result: 21 to 19

Results:
-----
ID_Improved          60.00%
```

We observe that it only has a slight edge over the Open Move heuristic.

Aggressive

Formula: `blank spaces - opponent moves`

This strategy was designed as direct counter to the Open strategy. Whereas the Open Move player plays defensively by favoring board states with more open moves, the Aggressive strategy actively tries to limit the opponent's moves.

The rationale for this strategy is that the horizon effect is particularly pronounced in this version of Isolation since moves are L-shaped 'knight' moves. Unlike regular isolation, each new board position has a completely different set of next possible positions from the previous. So while it may seem like moving to a position with a lot of open moves is a good idea, it may also be a trap where all the next possible moves will be the player's last. Thus, it may be more effective to go on the offensive.

Below is the result of its tournament run:

```
*****
Evaluating: Student Aggressive
*****

Playing Matches:
-----
Match 1: Student Aggressive vs Random      Result: 36 to 4
Match 2: Student Aggressive vs MM_Null     Result: 28 to 12
Match 3: Student Aggressive vs MM_Open     Result: 24 to 16
Match 4: Student Aggressive vs MM_Improved Result: 16 to 24
Match 5: Student Aggressive vs AB_Null     Result: 26 to 14
Match 6: Student Aggressive vs AB_Open     Result: 26 to 14
Match 7: Student Aggressive vs AB_Improved Result: 22 to 18

Results:
-----
Student Aggressive      63.57%
```

Overall, the results show that the strategy is slightly better than our baseline, at least in this tournament run. Due to the variance in results, it is difficult to conclusively say it performs better than our baseline. One good thing is that it seems to beat the Open Move heuristic by a noticeable margin. Another observation is that it frequently loses to MM_Improved but consistently beats AB_Improved in several simulations. One possible explanation is that the scorer is more accurate come mid to end game and Alpha-Beta optimization allows for much earlier access to those deeper nodes.

Balanced

Formula: $\text{player moves} * (\text{blank spaces} - \text{opponent moves})$

This scorer was conceived as a possible improvement over the Aggressive strategy by factoring in open moves similar to ID_Improved. Instead of addition, we use multiplication. The rationale is that the score is really two terms with different units and the player's open move may not be equal in value to the opponent's.

Below is the result of its tournament run:

```
*****
Evaluating: Student Balanced
*****

Playing Matches:
-----
Match 1: Student Balanced vs Random      Result: 34 to 6
Match 2: Student Balanced vs MM_Null      Result: 29 to 11
Match 3: Student Balanced vs MM_Open      Result: 17 to 23
Match 4: Student Balanced vs MM_Improved  Result: 20 to 20
Match 5: Student Balanced vs AB_Null      Result: 28 to 12
Match 6: Student Balanced vs AB_Open      Result: 24 to 16
Match 7: Student Balanced vs AB_Improved  Result: 25 to 15

Results:
-----
Student Balanced      63.21%
```

Unfortunately, it did not seem to do much better than the previous strategy, although in some simulations it can reach a winrate of about 67%. Again, we cannot conclusively say it does better than our baseline. We notice however, that it now frequently ties with MM_Improved. This could suggest that factoring in open moves has improved our player's mid game somewhat. Also worth noting that it is now losing some games to MM_Open, so our new factor seems to have introduced some weaknesses as well.

Monte Carlo

The last strategy uses a simple Monte Carlo (MC) tree search to estimate the viability of a move or board state. Starting from a board state we simulate a game with random moves to the end state and recording if its a win or loss. By running several simulations, we can obtain the ratio of wins over the number of simulations or in essence, a 'probability of winning' for that move. This somewhat counters the horizon effect and gives our player a 'blurry' glimpse of the future.

There were some hurdles getting this to work with the project setup, particularly with Iterative Deepening. It didn't seem particularly efficient to throw away perfectly good MC estimates so some tuning was necessary. After several trials, the optimal strategy seems to be applying a **time limit of 2ms** and a **max simulation count of 50** for evaluating each board state. The max simulation count doesn't matter as much as long as we set it to a sufficiently high value since the function tends to run out of time more often than not.

Below is the result of its tournament run:

```
*****
Evaluating: Student MCS
*****

Playing Matches:
-----
Match 1: Student MCS vs   Random      Result: 36 to 4
Match 2: Student MCS vs   MM_Null     Result: 34 to 6
Match 3: Student MCS vs   MM_Open     Result: 25 to 15
Match 4: Student MCS vs MM_Improved  Result: 27 to 13
Match 5: Student MCS vs   AB_Null     Result: 34 to 6
Match 6: Student MCS vs   AB_Open     Result: 27 to 13
Match 7: Student MCS vs AB_Improved  Result: 25 to 15

Results:
-----
Student MCS          74.29%
```

The result is a very encouraging **74.29%**. In other tournament runs, winrates were consistently above 70%. Moreover, it beats every other player by a decisive margin in each match.

All in all, this is a rather primitive implementation of Monte Carlo Tree Search. Additional improvements can be made by constrained sampling of movesets based on heuristics during simulation, for example. It's definitely possible to achieve much higher winrates with if MCTS was designed with more sophistication.

Note to the Reviewer

I should mention that the heuristic interface test will fail. This is because the 'player' parameter that is passed into the function is a string instead of a CustomPlayer and I needed the `time_left()` function for MCS.